

10 minutes to pandas

Intro to data structures

Essential basic functionality

IO tools (text, CSV, HDF5, ...)

PyArrow Functionality

Indexing and selecting data

MultIndex / advanced indexing

Copy-on-Write (CoW)

Merge, join, concatenate and
compare

Reshaping and pivot tables

Working with text data

Working with missing data

Duplicate Labels

Categorical data

Nullable integer data type

Nullable Boolean data type

Chart visualization

Table Visualization

User-Defined Functions (UDFs)

Essential basic functionality

Here we discuss a lot of the essential functionality common to the pandas data structures. To begin, let's create some example objects like we did in the [10 minutes to pandas](#) section:

```
In [1]: index = pd.date_range("1/1/2000", periods=8)
```

```
In [2]: s = pd.Series(np.random.randn(5), index=["a", "b", "c", "d", "e"])
```

```
In [3]: df = pd.DataFrame(np.random.randn(8, 3), index=index, columns=["A", "B", "C"])
```

Head and tail

To view a small sample of a Series or DataFrame object, use the `head()` and `tail()` methods. The default number of elements to display is five, but you may pass a custom number.

```
In [4]: long_series = pd.Series(np.random.randn(1000))
```

```
In [5]: long_series.head()
```

```
Out[5]:
```

```
0    -1.157892
1    -1.344312
2     0.844885
3     1.075770
4    -0.109050
dtype: float64
```

```
In [6]: long_series.tail(3)
```

```
Out[6]:
```

```
997    -0.289388
998    -1.020544
999     0.589993
dtype: float64
```

Attributes and underlying data

pandas objects have a number of attributes enabling you to access the metadata.

- **shape:** gives the axis dimensions of the object, consistent with ndarray
- **Axis labels**
 - **Series:** *index* (only axis)
 - **DataFrame:** *index* (rows) and *columns*

Note, **these attributes can be safely assigned to!**

```

In [7]: df[:2]
Out[7]:
           A          B          C
2000-01-01 -0.173215  0.119209 -1.044236
2000-01-02 -0.861849 -2.104569 -0.494929

In [8]: df.columns = [x.lower() for x in df.columns]

In [9]: df
Out[9]:
           a          b          c
2000-01-01 -0.173215  0.119209 -1.044236
2000-01-02 -0.861849 -2.104569 -0.494929
2000-01-03  1.071804  0.721555 -0.706771
2000-01-04 -1.039575  0.271860 -0.424972
2000-01-05  0.567020  0.276232 -1.087401
2000-01-06 -0.673690  0.113648 -1.478427
2000-01-07  0.524988  0.404705  0.577046
2000-01-08 -1.715002 -1.039268 -0.370647

```

pandas objects ([Index](#), [Series](#), [DataFrame](#)) can be thought of as containers for arrays, which hold the actual data and do the actual computation. For many types, the underlying array is a [numpy.ndarray](#). However, pandas and 3rd party libraries may *extend* NumPy's type system to add support for custom arrays (see [dtypes](#)).

To get the actual data inside a [Index](#) or [Series](#), use the [.array](#) property.

```
In [10]: s.array
Out[10]:
<NumpyExtensionArray>
[ 0.4691122999071863, -0.2828633443286633, -1.5090585031735124,
 -1.1356323710171934,  1.2121120250208506]
Length: 5, dtype: float64

In [11]: s.index.array
Out[11]:
<ArrowStringArray>
['a', 'b', 'c', 'd', 'e']
Length: 5, dtype: str
```

`array` will always be an `ExtensionArray`. The exact details of what an `ExtensionArray` is and why pandas uses them are a bit beyond the scope of this introduction. See [dtypes](#) for more.

If you know you need a NumPy array, use `to_numpy()` or `numpy.asarray()`.

```
In [12]: s.to_numpy()
Out[12]: array([ 0.4691, -0.2829, -1.5091, -1.1356,  1.2121])

In [13]: np.asarray(s)
Out[13]: array([ 0.4691, -0.2829, -1.5091, -1.1356,  1.2121])
```

When the Series or Index is backed by an `ExtensionArray`, `to_numpy()` may involve copying data and coercing values. See [dtypes](#) for more.

`to_numpy()` gives some control over the `dtype` of the resulting `numpy.ndarray`. For example, consider datetimes with timezones. NumPy doesn't have a dtype to represent timezone-aware datetimes, so there are two possibly useful representations:

1. An object-dtype `numpy.ndarray` with `Timestamp` objects, each with the correct `tz`.

2. A `datetime64[ns]` -dtype `numpy.ndarray`, where the values have been converted to UTC and the timezone discarded.

Timezones may be preserved with `dtype=object`:

```
In [14]: ser = pd.Series(pd.date_range("2000", periods=2, tz="CET"))

In [15]: ser.to_numpy(dtype=object)
Out[15]:
array([Timestamp('2000-01-01 00:00:00+0100', tz='CET'),
       Timestamp('2000-01-02 00:00:00+0100', tz='CET')], dtype=object)
```

Or thrown away with `dtype='datetime64[ns]'`:

```
In [16]: ser.to_numpy(dtype="datetime64[ns]")
Out[16]:
array(['1999-12-31T23:00:00.000000000', '2000-01-01T23:00:00.000000000'],
      dtype='datetime64[ns]')
```

Getting the “raw data” inside a `DataFrame` is possibly a bit more complex. When your `DataFrame` only has a single data type for all the columns, `DataFrame.to_numpy()` will return the underlying data:

```
In [17]: df.to_numpy()
Out[17]:
array([[ -0.1732,   0.1192,  -1.0442],
       [ -0.8618,  -2.1046,  -0.4949],
       [  1.0718,   0.7216,  -0.7068],
       [ -1.0396,   0.2719,  -0.425 ],
       [  0.567 ,   0.2762,  -1.0874],
       [ -0.6737,   0.1136,  -1.4784],
       [  0.525 ,   0.4047,   0.577 ],
       [ -1.715 ,  -1.0393,  -0.3706]])
```

If a DataFrame contains homogeneously-typed data, the ndarray can actually be modified in-place, and the changes will be reflected in the data structure. For heterogeneous data (e.g. some of the DataFrame's columns are not all the same dtype), this will not be the case. The values attribute itself, unlike the axis labels, cannot be assigned to.

Note

When working with heterogeneous data, the dtype of the resulting ndarray will be chosen to accommodate all of the data involved. For example, if strings are involved, the result will be of object dtype. If there are only floats and integers, the resulting array will be of float dtype.

In the past, pandas recommended `Series.values` or `DataFrame.values` for extracting the data from a Series or DataFrame. You'll still find references to these in old code bases and online. Going forward, we recommend avoiding `.values` and using `.array` or `.to_numpy()`. `.values` has the following drawbacks:

1. When your Series contains an [extension type](#), it's unclear whether `Series.values` returns a NumPy array or the extension array. `Series.array` will always return an `ExtensionArray`, and will never copy data. `Series.to_numpy()` will always return a NumPy array, potentially at the cost of copying / coercing values.
2. When your DataFrame contains a mixture of data types, `DataFrame.values` may involve copying data and coercing values to a common dtype, a relatively expensive operation. `DataFrame.to_numpy()`, being a method, makes it clearer that the returned NumPy array may not be a view on the same data in the DataFrame.

Accelerated operations

pandas has support for accelerating certain types of binary numerical and boolean operations using the `numexpr` library and the `bottleneck` libraries.

These libraries are especially useful when dealing with large data sets, and provide large speedups. `numexpr` uses smart chunking, caching, and multiple cores. `bottleneck` is a set of specialized cython routines that are especially fast when dealing with arrays that have `nan`.

You are highly encouraged to install both libraries. See the section [Recommended Dependencies](#) for more installation info.

These are both enabled to be used by default, you can control this by setting the options:

```
pd.set_option("compute.use_bottleneck", False)
pd.set_option("compute.use_numexpr", False)
```

Flexible binary operations

With binary operations between pandas data structures, there are two key points of interest:

- Broadcasting behavior between higher- (e.g. DataFrame) and lower-dimensional (e.g. Series) objects.
- Missing data in computations.

We will demonstrate how to manage these issues independently, though they can be handled simultaneously.

Matching / broadcasting behavior

DataFrame has the methods `add()`, `sub()`, `mul()`, `div()` and related functions `radd()`, `rsub()`, ... for carrying out binary operations. For broadcasting behavior, Series input is of primary interest. Using these functions, you can use to either match on the *index* or *columns* via the **axis** keyword:


```
In [18]: df = pd.DataFrame(
.....:     {
.....:         "one": pd.Series(np.random.randn(3), index=["a", "b", "c"]),
.....:         "two": pd.Series(np.random.randn(4), index=["a", "b", "c", "d"]),
.....:         "three": pd.Series(np.random.randn(3), index=["b", "c", "d"]),
.....:     }
.....: )
.....:
```

```
In [19]: df
```

```
Out[19]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [20]: row = df.iloc[1]
```

```
In [21]: column = df["two"]
```

```
In [22]: df.sub(row, axis="columns")
```

```
Out[22]:
```

	one	two	three
a	1.051928	-0.139606	NaN
b	0.000000	0.000000	0.000000
c	0.352192	-0.433754	1.277825
d	NaN	-1.632779	-0.562782

```
In [23]: df.sub(row, axis=1)
```

```
Out[23]:
```

	one	two	three
a	1.051928	-0.139606	NaN
b	0.000000	0.000000	0.000000
c	0.352192	-0.433754	1.277825
d	NaN	-1.632779	-0.562782

```
In [24]: df.sub(column, axis="index")
```

```
Out[24]:
```

	one	two	three
--	-----	-----	-------

```

a -0.377535  0.0      NaN
b -1.569069  0.0 -1.962513
c -0.783123  0.0 -0.250933
d          NaN  0.0 -0.892516

In [25]: df.sub(column, axis=0)
Out[25]:
      one  two      three
a -0.377535  0.0      NaN
b -1.569069  0.0 -1.962513
c -0.783123  0.0 -0.250933
d          NaN  0.0 -0.892516

```

Furthermore you can align a level of a MultiIndexed DataFrame with a Series.

```

In [26]: dfmi = df.copy()

In [27]: dfmi.index = pd.MultiIndex.from_tuples(
....:     [(1, "a"), (1, "b"), (1, "c"), (2, "a")], names=["first", "second"]
....: )
....:

In [28]: dfmi.sub(column, axis=0, level="second")
Out[28]:
      one  two      three
first second
1      a   -0.377535  0.000000      NaN
      b  -1.569069  0.000000 -1.962513
      c  -0.783123  0.000000 -0.250933
2      a          NaN -1.493173 -2.385688

```

Series and Index also support the `divmod()` builtin. This function takes the floor division and modulo operation at the same time returning a two-tuple of the same type as the left hand side. For example:

```
In [29]: s = pd.Series(np.arange(10))
```

```
In [30]: s
```

```
Out[30]:
```

```
0    0
1    1
2    2
3    3
4    4
5    5
6    6
7    7
8    8
9    9
```

```
dtype: int64
```

```
In [31]: div, rem = divmod(s, 3)
```

```
In [32]: div
```

```
Out[32]:
```

```
0    0
1    0
2    0
3    1
4    1
5    1
6    2
7    2
8    2
9    3
```

```
dtype: int64
```

```
In [33]: rem
```

```
Out[33]:
```

```
0    0
1    1
2    2
3    0
4    1
5    2
```

```
6    0
7    1
8    2
9    0
dtype: int64
```

```
In [34]: idx = pd.Index(np.arange(10))
```

```
In [35]: idx
```

```
Out[35]: Index([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype='int64')
```

```
In [36]: div, rem = divmod(idx, 3)
```

```
In [37]: div
```

```
Out[37]: Index([0, 0, 0, 1, 1, 1, 2, 2, 2, 3], dtype='int64')
```

```
In [38]: rem
```

```
Out[38]: Index([0, 1, 2, 0, 1, 2, 0, 1, 2, 0], dtype='int64')
```

We can also do elementwise `divmod()`:

```
In [39]: div, rem = divmod(s, [2, 2, 3, 3, 4, 4, 5, 5, 6, 6])
```

```
In [40]: div
```

```
Out[40]:
```

```
0    0
1    0
2    0
3    1
4    1
5    1
6    1
7    1
8    1
9    1
dtype: int64
```

```
In [41]: rem
```

```
Out[41]:
```

```
0    0
1    1
2    2
3    0
4    0
5    1
6    1
7    2
8    2
9    3
dtype: int64
```

Missing data / operations with fill values

In Series and DataFrame, the arithmetic functions have the option of inputting a *fill_value*, namely a value to substitute when at most one of the values at a location are missing. For example, when adding two DataFrame objects, you may wish to treat NaN as 0 unless both DataFrames are missing that value, in

which case the result will be NaN (you can later replace NaN with some other value using `fillna` if you wish).

```
In [42]: df2 = df.copy()
```

```
In [43]: df2.loc["a", "three"] = 1.0
```

```
In [44]: df
```

```
Out[44]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [45]: df2
```

```
Out[45]:
```

	one	two	three
a	1.394981	1.772517	1.000000
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [46]: df + df2
```

```
Out[46]:
```

	one	two	three
a	2.789963	3.545034	NaN
b	0.686107	3.824246	-0.100780
c	1.390491	2.956737	2.454870
d	NaN	0.558688	-1.226343

```
In [47]: df.add(df2, fill_value=0)
```

```
Out[47]:
```

	one	two	three
a	2.789963	3.545034	1.000000
b	0.686107	3.824246	-0.100780
c	1.390491	2.956737	2.454870
d	NaN	0.558688	-1.226343

Flexible comparisons

Series and DataFrame have the binary comparison methods `eq`, `ne`, `lt`, `gt`, `le`, and `ge` whose behavior is analogous to the binary arithmetic operations described above:

```
In [48]: df.gt(df2)
Out[48]:
```

	one	two	three
a	False	False	False
b	False	False	False
c	False	False	False
d	False	False	False

```
In [49]: df2.ne(df)
Out[49]:
```

	one	two	three
a	False	False	True
b	False	False	False
c	False	False	False
d	True	False	False

These operations produce a pandas object of the same type as the left-hand-side input that is of dtype `bool`. These `boolean` objects can be used in indexing operations, see the section on [Boolean indexing](#).

Boolean reductions

You can apply the reductions: `empty`, `any()`, `all()`.

```
In [50]: (df > 0).all()
Out[50]:
one      False
two       True
three    False
dtype: bool
```

```
In [51]: (df > 0).any()
Out[51]:
one       True
two       True
three     True
dtype: bool
```

You can reduce to a final boolean value.

```
In [52]: (df > 0).any().any()
Out[52]: np.True_
```

You can test if a pandas object is empty, via the `empty` property.

```
In [53]: df.empty
Out[53]: False

In [54]: pd.DataFrame(columns=list("ABC")).empty
Out[54]: True
```


Warning

Asserting the truthiness of a pandas object will raise an error, as the testing of the emptiness or values is ambiguous.

```
In [55]: if df:
.....:     print(True)
.....:

-----
ValueError                                Traceback (most recent call last)
Cell In[55], line 1
----> 1 if df:
      2     print(True)

File ~/work/pandas/pandas/pandas/core/generic.py:1501, in NDFrame.__bool__(self)
    1499     @final
    1500     def __bool__(self) -> NoReturn:
-> 1501         raise ValueError(
    1502             f"The truth value of a {type(self).__name__} is ambiguous. "
    1503             "Use a.empty, a.bool(), a.item(), a.any() or a.all()."
    1504         )

ValueError: The truth value of a DataFrame is ambiguous. Use a.empty, a.bool(), a.item()
```

```
In [56]: df and df2
```

```
-----  
ValueError
```

```
Traceback (most recent call last)
```

```
Cell In[56], line 1
```

```
----> 1 df and df2
```

```
File ~/work/pandas/pandas/pandas/core/generic.py:1501, in NDFrame.__bool__(self)
```

```
1499     @final
```

```
1500     def __bool__(self) -> NoReturn:
```

```
-> 1501         raise ValueError(
```

```
1502             f"The truth value of a {type(self).__name__} is ambiguous. "
```

```
1503             "Use a.empty, a.bool(), a.item(), a.any() or a.all()."
```

```
1504         )
```

```
ValueError: The truth value of a DataFrame is ambiguous. Use a.empty, a.bool(), a.item()
```

See [gotchas](#) for a more detailed discussion.

Comparing if objects are equivalent

Often you may find that there is more than one way to compute the same result. As a simple example, consider `df + df` and `df * 2`. To test that these two computations produce the same result, given the tools shown above, you might imagine using `(df + df == df * 2).all()`. But in fact, this expression is False:

```
In [57]: df + df == df * 2
Out[57]:
```

	one	two	three
a	True	True	False
b	True	True	True
c	True	True	True
d	False	True	True

```
In [58]: (df + df == df * 2).all()
Out[58]:
```

	one	two	three
one	False		
two		True	
three		False	

```
dtype: bool
```

Notice that the boolean DataFrame `df + df == df * 2` contains some False values! This is because NaNs do not compare as equals:

```
In [59]: np.nan == np.nan
Out[59]: False
```

So, NDFrames (such as Series and DataFrames) have an `equals()` method for testing equality, with NaNs in corresponding locations treated as equal.

```
In [60]: (df + df).equals(df * 2)
Out[60]: True
```

Note that the Series or DataFrame index needs to be in the same order for equality to be True:

```
In [61]: df1 = pd.DataFrame({"col": ["foo", 0, np.nan]})

In [62]: df2 = pd.DataFrame({"col": [np.nan, 0, "foo"]}, index=[2, 1, 0])

In [63]: df1.equals(df2)
Out[63]: False

In [64]: df1.equals(df2.sort_index())
Out[64]: True
```

Comparing array-like objects

You can conveniently perform element-wise comparisons when comparing a pandas data structure with a scalar value:

```
In [65]: pd.Series(["foo", "bar", "baz"]) == "foo"
Out[65]:
0      True
1     False
2     False
dtype: bool

In [66]: pd.Index(["foo", "bar", "baz"]) == "foo"
Out[66]: array([ True, False, False])
```

pandas also handles element-wise comparisons between different array-like objects of the same length:

```
In [67]: pd.Series(["foo", "bar", "baz"]) == pd.Index(["foo", "bar", "qux"])
Out[67]:
0      True
1      True
2     False
dtype: bool

In [68]: pd.Series(["foo", "bar", "baz"]) == np.array(["foo", "bar", "qux"])
Out[68]:
0      True
1      True
2     False
dtype: bool
```

Trying to compare `Index` or `Series` objects of different lengths will raise a `ValueError`:

```

In [69]: pd.Series(['foo', 'bar', 'baz']) == pd.Series(['foo', 'bar'])
-----
ValueError                                Traceback (most recent call last)
Cell In[69], line 1
----> 1 pd.Series(['foo', 'bar', 'baz']) == pd.Series(['foo', 'bar'])

File ~/work/pandas/pandas/pandas/core/ops/common.py:85, in _unpack_zerodim_and_defer.<locals>.ne
    82     other = sanitize_array(other, None)
    83     other = ensure_wrapped_if_datetimelike(other)
----> 85 return method(self, other)

File ~/work/pandas/pandas/pandas/core/arraylike.py:42, in OpsMixin.__eq__(self, other)
    40 @unpack_zerodim_and_defer("__eq__")
    41 def __eq__(self, other):
----> 42     return self._cmp_method(other, operator.eq)

File ~/work/pandas/pandas/pandas/core/series.py:6730, in Series._cmp_method(self, other, op)
    6727 res_name = ops.get_op_result_name(self, other)
    6729 if isinstance(other, Series) and not self._indexed_same(other):
-> 6730     raise ValueError("Can only compare identically-labeled Series objects")
    6732 lvalues = self._values
    6733 rvalues = extract_array(other, extract_numpy=True, extract_range=True)

ValueError: Can only compare identically-labeled Series objects

In [70]: pd.Series(['foo', 'bar', 'baz']) == pd.Series(['foo'])
-----
ValueError                                Traceback (most recent call last)
Cell In[70], line 1
----> 1 pd.Series(['foo', 'bar', 'baz']) == pd.Series(['foo'])

File ~/work/pandas/pandas/pandas/core/ops/common.py:85, in _unpack_zerodim_and_defer.<locals>.ne
    82     other = sanitize_array(other, None)
    83     other = ensure_wrapped_if_datetimelike(other)
----> 85 return method(self, other)

File ~/work/pandas/pandas/pandas/core/arraylike.py:42, in OpsMixin.__eq__(self, other)
    40 @unpack_zerodim_and_defer("__eq__")
    41 def __eq__(self, other):
----> 42     return self._cmp_method(other, operator.eq)

```

```
File ~/work/pandas/pandas/pandas/core/series.py:6730, in Series._cmp_method(self, other, op)
    6727 res_name = ops.get_op_result_name(self, other)
    6729 if isinstance(other, Series) and not self._indexed_same(other):
-> 6730     raise ValueError("Can only compare identically-labeled Series objects")
    6732 lvalues = self._values
    6733 rvalues = extract_array(other, extract_numpy=True, extract_range=True)
```

```
ValueError: Can only compare identically-labeled Series objects
```

Combining overlapping data sets

A problem occasionally arising is the combination of two similar data sets where values in one are preferred over the other. An example would be two data series representing a particular economic indicator where one is considered to be of "higher quality". However, the lower quality series might extend further back in history or have more complete data coverage. As such, we would like to combine two DataFrame objects where missing values in one DataFrame are conditionally filled with like-labeled values from the other DataFrame. The function implementing this operation is `combine_first()`, which we illustrate:

```
In [71]: df1 = pd.DataFrame(
.....:     {"A": [1.0, np.nan, 3.0, 5.0, np.nan], "B": [np.nan, 2.0, 3.0, np.nan, 6.0]}
.....: )
.....:
```

```
In [72]: df2 = pd.DataFrame(
.....:     {
.....:         "A": [5.0, 2.0, 4.0, np.nan, 3.0, 7.0],
.....:         "B": [np.nan, np.nan, 3.0, 4.0, 6.0, 8.0],
.....:     }
.....: )
.....:
```

```
In [73]: df1
```

```
Out[73]:
```

	A	B
0	1.0	NaN
1	NaN	2.0
2	3.0	3.0
3	5.0	NaN
4	NaN	6.0

```
In [74]: df2
```

```
Out[74]:
```

	A	B
0	5.0	NaN
1	2.0	NaN
2	4.0	3.0
3	NaN	4.0
4	3.0	6.0
5	7.0	8.0

```
In [75]: df1.combine_first(df2)
```

```
Out[75]:
```

	A	B
0	1.0	NaN
1	2.0	2.0
2	3.0	3.0
3	5.0	4.0


```
4  3.0  6.0
5  7.0  8.0
```

General DataFrame combine

The `combine_first()` method above calls the more general `DataFrame.combine()`. This method takes another DataFrame and a combiner function, aligns the input DataFrame and then passes the combiner function pairs of Series (i.e., columns whose names are the same).

So, for instance, to reproduce `combine_first()` as above:

```
In [76]: def combiner(x, y):
....:     return np.where(pd.isna(x), y, x)
....:
```

```
In [77]: df1.combine(df2, combiner)
```

```
Out[77]:
```

```
   A  B
0  1.0 NaN
1  2.0 2.0
2  3.0 3.0
3  5.0 4.0
4  3.0 6.0
5  7.0 8.0
```

Descriptive statistics

There exists a large number of methods for computing descriptive statistics and other related operations on [Series](#), [DataFrame](#). Most of these are aggregations (hence producing a lower-dimensional result) like

`sum()`, `mean()`, and `quantile()`, but some of them, like `cumsum()` and `cumprod()`, produce an object of the same size. Generally speaking, these methods take an **axis** argument, just like `ndarray.{sum, std, ...}`, but the axis can be specified by name or integer:

- **Series:** no axis argument needed
- **DataFrame:** "index" (axis=0, default), "columns" (axis=1)

For example:

```
In [78]: df
Out[78]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [79]: df.mean(axis=0)
Out[79]:
```

one	0.811094
two	1.360588
three	0.187958

```
dtype: float64

In [80]: df.mean(axis=1)
Out[80]:
```

a	1.583749
b	0.734929
c	1.133683
d	-0.166914

```
dtype: float64
```

All such methods have a `skipna` option signaling whether to exclude missing data (`True` by default):

```
In [81]: df.sum(axis=0, skipna=False)
```

```
Out[81]:
```

```
one      NaN
two      5.442353
three    NaN
dtype: float64
```

```
In [82]: df.sum(axis=1, skipna=True)
```

```
Out[82]:
```

```
a      3.167498
b      2.204786
c      3.401050
d     -0.333828
dtype: float64
```

Combined with the broadcasting / arithmetic behavior, one can describe various statistical procedures, like standardization (rendering data zero mean and standard deviation of 1), very concisely:

```
In [83]: ts_stand = (df - df.mean()) / df.std()
```

```
In [84]: ts_stand.std()
```

```
Out[84]:
```

```
one      1.0
two      1.0
three    1.0
dtype: float64
```

```
In [85]: xs_stand = df.sub(df.mean(axis=1), axis=0).div(df.std(axis=1), axis=0)
```

```
In [86]: xs_stand.std(axis=1)
```

```
Out[86]:
```

```
a      1.0
b      1.0
c      1.0
d      1.0
dtype: float64
```

Note that methods like `cumsum()` and `cumprod()` preserve the location of `NaN` values. This is somewhat different from `expanding()` and `rolling()` since `NaN` behavior is furthermore dictated by a `min_periods` parameter.

```
In [87]: df.cumsum()
Out[87]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	1.738035	3.684640	-0.050390
c	2.433281	5.163008	1.177045
d	NaN	5.442353	0.563873

Here is a quick reference summary table of common functions. Each also takes an optional `level` parameter which applies only if the object has a [hierarchical index](#).

Function	Description
<code>count</code>	Number of non-NA observations
<code>sum</code>	Sum of values
<code>mean</code>	Mean of values
<code>median</code>	Arithmetic median of values
<code>min</code>	Minimum
<code>max</code>	Maximum
<code>mode</code>	Mode

Function	Description
<code>abs</code>	Absolute Value
<code>prod</code>	Product of values
<code>std</code>	Bessel-corrected sample standard deviation
<code>var</code>	Unbiased variance
<code>sem</code>	Standard error of the mean
<code>skew</code>	Sample skewness (3rd moment)
<code>kurt</code>	Sample kurtosis (4th moment)
<code>quantile</code>	Sample quantile (value at %)
<code>cumsum</code>	Cumulative sum
<code>cumprod</code>	Cumulative product
<code>cummax</code>	Cumulative maximum
<code>cummin</code>	Cumulative minimum

Note that by chance some NumPy methods, like `mean`, `std`, and `sum`, will exclude NAs on Series input by default:

```
In [88]: np.mean(df["one"])
Out[88]: np.float64(0.8110935116651192)

In [89]: np.mean(df["one"].to_numpy())
Out[89]: np.float64(nan)
```

`Series.nunique()` will return the number of unique non-NA values in a Series:

```
In [90]: series = pd.Series(np.random.randn(500))

In [91]: series[20:500] = np.nan

In [92]: series[10:20] = 5

In [93]: series.nunique()
Out[93]: 11
```

Summarizing data: describe

There is a convenient `describe()` function which computes a variety of summary statistics about a Series or the columns of a DataFrame (excluding NAs of course):

```

In [94]: series = pd.Series(np.random.randn(1000))

In [95]: series[::2] = np.nan

In [96]: series.describe()
Out[96]:
count      500.000000
mean       -0.021292
std         1.015906
min        -2.683763
25%        -0.699070
50%        -0.069718
75%         0.714483
max         3.160915
dtype: float64

In [97]: frame = pd.DataFrame(np.random.randn(1000, 5), columns=["a", "b", "c", "d", "e"])

In [98]: frame.iloc[::2] = np.nan

In [99]: frame.describe()
Out[99]:

```

	a	b	c	d	e
count	500.000000	500.000000	500.000000	500.000000	500.000000
mean	0.033387	0.030045	-0.043719	-0.051686	0.005979
std	1.017152	0.978743	1.025270	1.015988	1.006695
min	-3.000951	-2.637901	-3.303099	-3.159200	-3.188821
25%	-0.647623	-0.576449	-0.712369	-0.691338	-0.691115
50%	0.047578	-0.021499	-0.023888	-0.032652	-0.025363
75%	0.729907	0.775880	0.618896	0.670047	0.649748
max	2.740139	2.752332	3.004229	2.728702	3.240991

You can select specific percentiles to include in the output:

```
In [100]: series.describe(percentiles=[0.05, 0.25, 0.75, 0.95])
Out[100]:
count      500.000000
mean       -0.021292
std         1.015906
min        -2.683763
5%         -1.645423
25%        -0.699070
75%         0.714483
95%         1.711409
max         3.160915
dtype: float64
```

By default, the median is always included.

For a non-numerical Series object, `describe()` will give a simple summary of the number of unique values and most frequently occurring values:

```
In [101]: s = pd.Series(["a", "a", "b", "b", "a", "a", np.nan, "c", "d", "a"])

In [102]: s.describe()
Out[102]:
count      9
unique      4
top         a
freq        5
dtype: object
```

Note that on a mixed-type DataFrame object, `describe()` will restrict the summary to include only numerical columns or, if none are, only categorical columns:


```
In [103]: frame = pd.DataFrame({"a": ["Yes", "Yes", "No", "No"], "b": range(4)})
```

```
In [104]: frame.describe()
```

```
Out[104]:
```

	b
count	4.000000
mean	1.500000
std	1.290994
min	0.000000
25%	0.750000
50%	1.500000
75%	2.250000
max	3.000000

This behavior can be controlled by providing a list of types as `include`/`exclude` arguments. The special value `all` can also be used:

```
In [105]: frame.describe(include=["str"])
```

```
Out[105]:
```

	a
count	4
unique	2
top	Yes
freq	2

```
In [106]: frame.describe(include=["number"])
```

```
Out[106]:
```

	b
count	4.000000
mean	1.500000
std	1.290994
min	0.000000
25%	0.750000
50%	1.500000
75%	2.250000
max	3.000000

```
In [107]: frame.describe(include="all")
```

```
Out[107]:
```

	a	b
count	4	4.000000
unique	2	NaN
top	Yes	NaN
freq	2	NaN
mean	NaN	1.500000
std	NaN	1.290994
min	NaN	0.000000
25%	NaN	0.750000
50%	NaN	1.500000
75%	NaN	2.250000
max	NaN	3.000000

That feature relies on [select_dtypes](#). Refer to there for details about accepted inputs.

Index of min/max values

The `idxmin()` and `idxmax()` functions on Series and DataFrame compute the index labels with the minimum and maximum corresponding values:

```
In [108]: s1 = pd.Series(np.random.randn(5))
```

```
In [109]: s1
```

```
Out[109]:
```

```
0    1.118076
1   -0.352051
2   -1.242883
3   -1.277155
4   -0.641184
dtype: float64
```

```
In [110]: s1.idxmin(), s1.idxmax()
```

```
Out[110]: (3, 0)
```

```
In [111]: df1 = pd.DataFrame(np.random.randn(5, 3), columns=["A", "B", "C"])
```

```
In [112]: df1
```

```
Out[112]:
```

	A	B	C
0	-0.327863	-0.946180	-0.137570
1	-0.186235	-0.257213	-0.486567
2	-0.507027	-0.871259	-0.111110
3	2.000339	-2.430505	0.089759
4	-0.321434	-0.033695	0.096271

```
In [113]: df1.idxmin(axis=0)
```

```
Out[113]:
```

```
A    2
B    3
C    1
dtype: int64
```

```
In [114]: df1.idxmax(axis=1)
```

```
Out[114]:
```

```
0    C
1    A
2    C
3    A
```

```
4    C
dtype: str
```

When there are multiple rows (or columns) matching the minimum or maximum value, `idxmin()` and `idxmax()` return the first matching index:

```
In [115]: df3 = pd.DataFrame([2, 1, 1, 3, np.nan], columns=["A"], index=list("edcba"))
```

```
In [116]: df3
```

```
Out[116]:
```

```
     A
e  2.0
d  1.0
c  1.0
b  3.0
a  NaN
```

```
In [117]: df3["A"].idxmin()
```

```
Out[117]: 'd'
```

Note

`idxmin` and `idxmax` are called `argmin` and `argmax` in NumPy.

Value counts (histogramming) / mode

The `value_counts()` Series method computes a histogram of a 1D array of values. It can also be used as a function on regular arrays:

```

In [118]: data = np.random.randint(0, 7, size=50)

In [119]: data
Out[119]:
array([6, 6, 2, 3, 5, 3, 2, 5, 4, 5, 4, 3, 4, 5, 0, 2, 0, 4, 2, 0, 3, 2,
        2, 5, 6, 5, 3, 4, 6, 4, 3, 5, 6, 4, 3, 6, 2, 6, 6, 2, 3, 4, 2, 1,
        6, 2, 6, 1, 5, 4])

In [120]: s = pd.Series(data)

In [121]: s.value_counts()
Out[121]:
6    10
2    10
4     9
3     8
5     8
0     3
1     2
Name: count, dtype: int64

```

The `value_counts()` method can be used to count combinations across multiple columns. By default all columns are used but a subset can be selected using the `subset` argument.

```

In [122]: data = {"a": [1, 2, 3, 4], "b": ["x", "x", "y", "y"]}

In [123]: frame = pd.DataFrame(data)

In [124]: frame.value_counts()
Out[124]:
a  b
1  x    1
2  x    1
3  y    1
4  y    1
Name: count, dtype: int64

```

Similarly, you can get the most frequently occurring value(s), i.e. the mode, of the values in a Series or DataFrame:

```
In [125]: s5 = pd.Series([1, 1, 3, 3, 3, 5, 5, 7, 7, 7])

In [126]: s5.mode()
Out[126]:
0      3
1      7
dtype: int64

In [127]: df5 = pd.DataFrame(
.....:     {
.....:         "A": np.random.randint(0, 7, size=50),
.....:         "B": np.random.randint(-10, 15, size=50),
.....:     }
.....: )
.....:

In [128]: df5.mode()
Out[128]:
   A  B
0  1.0 -9
1  NaN 10
2  NaN 13
```

Discretization and quantiling

Continuous values can be discretized using the `cut()` (bins based on values) and `qcut()` (bins based on sample quantiles) functions:

```

In [129]: arr = np.random.randn(20)

In [130]: factor = pd.cut(arr, 4)

In [131]: factor
Out[131]:
[(-0.251, 0.464], (-0.968, -0.251], (0.464, 1.179], (-0.251, 0.464], (-0.968, -0.251], ..., (-0.
Length: 20
Categories (4, interval[float64, right]): [(-0.968, -0.251] < (-0.251, 0.464] < (0.464, 1.179] <
(1.179, 1.893]]

In [132]: factor = pd.cut(arr, [-5, -1, 0, 1, 5])

In [133]: factor
Out[133]:
[(0, 1], (-1, 0], (0, 1], (0, 1], (-1, 0], ..., (-1, 0], (-1, 0], (-1, 0], (-1, 0], (-1, 0]]
Length: 20
Categories (4, interval[int64, right]): [(-5, -1] < (-1, 0] < (0, 1] < (1, 5]]

```

[`qcut\(\)`](#) computes sample quantiles. For example, we could slice up some normally distributed data into equal-size quartiles like so:

```

In [134]: arr = np.random.randn(30)

In [135]: factor = pd.qcut(arr, [0, 0.25, 0.5, 0.75, 1])

In [136]: factor
Out[136]:
[(0.569, 1.184], (-2.278, -0.301], (-2.278, -0.301], (0.569, 1.184], (0.569, 1.184], ..., (-0.301, 0.569]
Length: 30
Categories (4, interval[float64, right]): [(-2.278, -0.301] < (-0.301, 0.569] < (0.569, 1.184] <
(1.184, 2.346]]

```


We can also pass infinite values to define the bins:

```
In [137]: arr = np.random.randn(20)

In [138]: factor = pd.cut(arr, [-np.inf, 0, np.inf])

In [139]: factor
Out[139]:
[(-inf, 0.0], (0.0, inf], (0.0, inf], (-inf, 0.0], (-inf, 0.0], ..., (-inf, 0.0], (-inf, 0.0], (
Length: 20
Categories (2, interval[float64, right]): [(-inf, 0.0] < (0.0, inf]]
```

Function application

To apply your own or another library's functions to pandas objects, you should be aware of the three methods below. The appropriate method to use depends on whether your function expects to operate on an entire `DataFrame` or `Series`, row- or column-wise, or elementwise.

1. [Tablewise Function Application](#): `pipe()`
2. [Row or Column-wise Function Application](#): `apply()`
3. [Aggregation API](#): `agg()` and `transform()`
4. [Applying Elementwise Functions](#): `map()`

Tablewise function application

`DataFrames` and `Series` can be passed into functions. However, if the function needs to be called in a chain, consider using the `pipe()` method.

First some setup:

```
In [140]: def extract_city_name(df):
.....:     """
.....:     Chicago, IL -> Chicago for city_name column
.....:     """
.....:     df["city_name"] = df["city_and_code"].str.split(",").str.get(0)
.....:     return df
.....:

In [141]: def add_country_name(df, country_name=None):
.....:     """
.....:     Chicago -> Chicago-US for city_name column
.....:     """
.....:     col = "city_name"
.....:     df["city_and_country"] = df[col] + country_name
.....:     return df
.....:

In [142]: df_p = pd.DataFrame({"city_and_code": ["Chicago, IL"]})
```

`extract_city_name` and `add_country_name` are functions taking and returning `DataFrames`.

Now compare the following:

```
In [143]: add_country_name(extract_city_name(df_p), country_name="US")
Out[143]:
  city_and_code city_name city_and_country
0  Chicago, IL   Chicago      ChicagoUS
```

Is equivalent to:

```
In [144]: df_p.pipe(extract_city_name).pipe(add_country_name, country_name="US")
Out[144]:
```

	city_and_code	city_name	city_and_country
0	Chicago, IL	Chicago	ChicagoUS

pandas encourages the second style, which is known as method chaining. `pipe` makes it easy to use your own or another library's functions in method chains, alongside pandas' methods.

In the example above, the functions `extract_city_name` and `add_country_name` each expected a `DataFrame` as the first positional argument. What if the function you wish to apply takes its data as, say, the second argument? In this case, provide `pipe` with a tuple of `(callable, data_keyword)`. `.pipe` will route the `DataFrame` to the argument specified in the tuple.

For example, we can fit a regression using statsmodels. Their API expects a formula first and a `DataFrame` as the second argument, `data`. We pass in the function, keyword pair `(sm.ols, 'data')` to `pipe`:

```
In [147]: import statsmodels.formula.api as sm
```

```
In [148]: bb = pd.read_csv("data/baseball.csv", index_col="id")
```

```
In [149]: (
.....:     bb.query("h > 0")
.....:     .assign(ln_h=lambda df: np.log(df.h))
.....:     .pipe((sm.ols, "data"), "hr ~ ln_h + year + g + C(lg)")
.....:     .fit()
.....:     .summary()
.....: )
.....:
```

```
Out[149]:
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```
=====
Dep. Variable:          hr      R-squared:            0.685
Model:                  OLS      Adj. R-squared:       0.665
Method:                 Least Squares      F-statistic:      34.28
Date:                  Tue, 22 Nov 2022      Prob (F-statistic):  3.48e-15
Time:                  05:34:17      Log-Likelihood:     -205.92
No. Observations:      68      AIC:                421.8
Df Residuals:          63      BIC:                432.9
Df Model:               4
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-8484.7720	4664.146	-1.819	0.074	-1.78e+04	835.780
C(lg)[T.NL]	-2.2736	1.325	-1.716	0.091	-4.922	0.375
ln_h	-1.3542	0.875	-1.547	0.127	-3.103	0.395
year	4.2277	2.324	1.819	0.074	-0.417	8.872
g	0.1841	0.029	6.258	0.000	0.125	0.243

```
=====
Omnibus:                10.875      Durbin-Watson:           1.999
Prob(Omnibus):          0.004      Jarque-Bera (JB):        17.298
Skew:                   0.537      Prob(JB):                0.000175
Kurtosis:               5.225      Cond. No.:               1.49e+07
=====
```

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 1.49e+07. This might indicate that there are  
strong multicollinearity or other numerical problems.  
"""
```

The pipe method is inspired by unix pipes and more recently [dplyr](#) and [magrittr](#), which have introduced the popular `(%>%)` (read pipe) operator for [R](#). The implementation of `pipe` here is quite clean and feels right at home in Python. We encourage you to view the source code of `pipe()`.

Row or column-wise function application

Arbitrary functions can be applied along the axes of a DataFrame using the `apply()` method, which, like the descriptive statistics methods, takes an optional `axis` argument:

```
In [145]: df.apply(lambda x: np.mean(x))
```

```
Out[145]:
```

```
one      0.811094
two      1.360588
three    0.187958
dtype: float64
```

```
In [146]: df.apply(lambda x: np.mean(x), axis=1)
```

```
Out[146]:
```

```
a      1.583749
b      0.734929
c      1.133683
d     -0.166914
dtype: float64
```

```
In [147]: df.apply(lambda x: x.max() - x.min())
```

```
Out[147]:
```

```
one      1.051928
two      1.632779
three    1.840607
dtype: float64
```

```
In [148]: df.apply(np.cumsum)
```

```
Out[148]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	1.738035	3.684640	-0.050390
c	2.433281	5.163008	1.177045
d	NaN	5.442353	0.563873

```
In [149]: df.apply(np.exp)
```

```
Out[149]:
```

	one	two	three
a	4.034899	5.885648	NaN
b	1.409244	6.767440	0.950858
c	2.004201	4.385785	3.412466
d	NaN	1.322262	0.541630

The `apply()` method will also dispatch on a string method name.

```
In [150]: df.apply("mean")
Out[150]:
one      0.811094
two      1.360588
three    0.187958
dtype: float64

In [151]: df.apply("mean", axis=1)
Out[151]:
a      1.583749
b      0.734929
c      1.133683
d     -0.166914
dtype: float64
```

The return type of the function passed to `apply()` affects the type of the final output from `DataFrame.apply` for the default behaviour:

- If the applied function returns a `Series`, the final output is a `DataFrame`. The columns match the index of the `Series` returned by the applied function.
- If the applied function returns any other type, the final output is a `Series`.

This default behaviour can be overridden using the `result_type`, which accepts three options: `reduce`, `broadcast`, and `expand`. These will determine how list-like return values expand (or not) to a `DataFrame`.

`apply()` combined with some cleverness can be used to answer many questions about a data set. For example, suppose we wanted to extract the date where the maximum value for each column occurred:

```

In [152]: tsdf = pd.DataFrame(
.....:     np.random.randn(1000, 3),
.....:     columns=["A", "B", "C"],
.....:     index=pd.date_range("1/1/2000", periods=1000),
.....: )
.....:

In [153]: tsdf.apply(lambda x: x.idxmax())
Out[153]:
A    2000-08-06
B    2001-01-18
C    2001-07-18
dtype: datetime64[us]

```

You may also pass additional arguments and keyword arguments to the `apply()` method.

```

In [154]: def subtract_and_divide(x, sub, divide=1):
.....:     return (x - sub) / divide
.....:

In [155]: df_udf = pd.DataFrame(np.ones((2, 2)))

In [156]: df_udf.apply(subtract_and_divide, args=(5,), divide=3)
Out[156]:
      0      1
0 -1.333333 -1.333333
1 -1.333333 -1.333333

```

Another useful feature is the ability to pass Series methods to carry out some Series operation on each column or row:


```
In [157]: tsdf = pd.DataFrame(
.....:     np.random.randn(10, 3),
.....:     columns=["A", "B", "C"],
.....:     index=pd.date_range("1/1/2000", periods=10),
.....: )
.....:
```

```
In [158]: tsdf.iloc[3:7] = np.nan
```

```
In [159]: tsdf
```

```
Out[159]:
```

	A	B	C
2000-01-01	-0.158131	-0.232466	0.321604
2000-01-02	-1.810340	-3.105758	0.433834
2000-01-03	-1.209847	-1.156793	-0.136794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	-0.653602	0.178875	1.008298
2000-01-09	1.007996	0.462824	0.254472
2000-01-10	0.307473	0.600337	1.643950

```
In [160]: tsdf.apply(pd.Series.interpolate)
```

```
Out[160]:
```

	A	B	C
2000-01-01	-0.158131	-0.232466	0.321604
2000-01-02	-1.810340	-3.105758	0.433834
2000-01-03	-1.209847	-1.156793	-0.136794
2000-01-04	-1.098598	-0.889659	0.092225
2000-01-05	-0.987349	-0.622526	0.321243
2000-01-06	-0.876100	-0.355392	0.550262
2000-01-07	-0.764851	-0.088259	0.779280
2000-01-08	-0.653602	0.178875	1.008298
2000-01-09	1.007996	0.462824	0.254472
2000-01-10	0.307473	0.600337	1.643950

Finally, `apply()` takes an argument `raw` which is `False` by default, which converts each row or column into a Series before applying the function. When set to `True`, the passed function will instead receive an ndarray object, which has positive performance implications if you do not need the indexing functionality.

Aggregation API

The aggregation API allows one to express possibly multiple aggregation operations in a single concise way. This API is similar across pandas objects, see [groupby API](#), the [window API](#), and the [resample API](#). The entry point for aggregation is `DataFrame.agg()`, or the alias `DataFrame.agg()`.

We will use a similar starting frame from above:

```
In [161]: tsdf = pd.DataFrame(
.....:     np.random.randn(10, 3),
.....:     columns=["A", "B", "C"],
.....:     index=pd.date_range("1/1/2000", periods=10),
.....: )
.....:
```

```
In [162]: tsdf.iloc[3:7] = np.nan
```

```
In [163]: tsdf
```

```
Out[163]:
```

	A	B	C
2000-01-01	1.257606	1.004194	0.167574
2000-01-02	-0.749892	0.288112	-0.757304
2000-01-03	-0.207550	-0.298599	0.116018
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.814347	-0.257623	0.869226
2000-01-09	-0.250663	-1.206601	0.896839
2000-01-10	2.169758	-1.333363	0.283157

Using a single function is equivalent to `apply()`. You can also pass named methods as strings. These will return a `Series` of the aggregated output:

```
In [164]: tsdf.agg(lambda x: np.sum(x))
Out[164]:
A    3.033606
B   -1.803879
C    1.575510
dtype: float64

In [165]: tsdf.agg("sum")
Out[165]:
A    3.033606
B   -1.803879
C    1.575510
dtype: float64

# these are equivalent to a ``.sum()`` because we are aggregating
# on a single function
In [166]: tsdf.sum()
Out[166]:
A    3.033606
B   -1.803879
C    1.575510
dtype: float64
```

Single aggregations on a `Series` this will return a scalar value:

```
In [167]: tsdf["A"].agg("sum")
Out[167]: np.float64(3.033606102414146)
```

Aggregating with multiple functions

You can pass multiple aggregation arguments as a list. The results of each of the passed functions will be a row in the resulting `DataFrame`. These are naturally named from the aggregation function.

```
In [168]: tsdf.agg(["sum"])
Out[168]:
```

	A	B	C
sum	3.033606	-1.803879	1.57551

Multiple functions yield multiple rows:

```
In [169]: tsdf.agg(["sum", "mean"])
Out[169]:
```

	A	B	C
sum	3.033606	-1.803879	1.575510
mean	0.505601	-0.300647	0.262585

On a `Series`, multiple functions return a `Series`, indexed by the function names:

```
In [170]: tsdf["A"].agg(["sum", "mean"])
Out[170]:
```

sum	3.033606
mean	0.505601

Name: A, dtype: float64

Passing a `lambda` function will yield a `<lambda>` named row:

```
In [171]: tsdf["A"].agg(["sum", lambda x: x.mean()])
Out[171]:
```

sum	3.033606
<lambda>	0.505601

Name: A, dtype: float64

Passing a named function will yield that name for the row:

```
In [172]: def mymean(x):
.....:     return x.mean()
.....:

In [173]: tsdf["A"].agg(["sum", mymean])
Out[173]:
sum      3.033606
mymean    0.505601
Name: A, dtype: float64
```

Aggregating with a dict

Passing a dictionary of column names to a scalar or a list of scalars, to `DataFrame.agg` allows you to customize which functions are applied to which columns. Note that the results are not in any particular order, you can use an `OrderedDict` instead to guarantee ordering.

```
In [174]: tsdf.agg({"A": "mean", "B": "sum"})
Out[174]:
A      0.505601
B     -1.803879
dtype: float64
```

Passing a list-like will generate a `DataFrame` output. You will get a matrix-like output of all of the aggregators. The output will consist of all unique functions. Those that are not noted for a particular column will be `NaN`:

```
In [175]: tsdf.agg({"A": ["mean", "min"], "B": "sum"})
```

```
Out[175]:
```

	A	B
mean	0.505601	NaN
min	-0.749892	NaN
sum	NaN	-1.803879

Custom describe

With `.agg()` it is possible to easily create a custom describe function, similar to the built in [describe function](#).

```
In [176]: from functools import partial
```

```
In [177]: q_25 = partial(pd.Series.quantile, q=0.25)
```

```
In [178]: q_25.__name__ = "25%"
```

```
In [179]: q_75 = partial(pd.Series.quantile, q=0.75)
```

```
In [180]: q_75.__name__ = "75%"
```

```
In [181]: tsdf.agg(["count", "mean", "std", "min", q_25, "median", q_75, "max"])
```

```
Out[181]:
```

	A	B	C
count	6.000000	6.000000	6.000000
mean	0.505601	-0.300647	0.262585
std	1.103362	0.887508	0.606860
min	-0.749892	-1.333363	-0.757304
25%	-0.239885	-0.979600	0.128907
median	0.303398	-0.278111	0.225365
75%	1.146791	0.151678	0.722709
max	2.169758	1.004194	0.896839

Transform API

The `transform()` method returns an object that is indexed the same (same size) as the original. This API allows you to provide *multiple* operations at the same time rather than one-by-one. Its API is quite similar to the `.agg` API.

We create a frame similar to the one used in the above sections.

```
In [182]: tsdf = pd.DataFrame(
.....:     np.random.randn(10, 3),
.....:     columns=["A", "B", "C"],
.....:     index=pd.date_range("1/1/2000", periods=10),
.....: )
.....:
```

```
In [183]: tsdf.iloc[3:7] = np.nan
```

```
In [184]: tsdf
```

```
Out[184]:
```

	A	B	C
2000-01-01	-0.428759	-0.864890	-0.675341
2000-01-02	-0.168731	1.338144	-1.279321
2000-01-03	-1.621034	0.438107	0.903794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	-1.240447	-0.201052
2000-01-09	-0.157795	0.791197	-1.144209
2000-01-10	-0.030876	0.371900	0.061932

Transform the entire frame. `.transform()` allows input functions as: a NumPy function, a string function name or a user defined function.


```
In [185]: tsdf.transform(np.abs)
```

```
Out[185]:
```

	A	B	C
2000-01-01	0.428759	0.864890	0.675341
2000-01-02	0.168731	1.338144	1.279321
2000-01-03	1.621034	0.438107	0.903794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	1.240447	0.201052
2000-01-09	0.157795	0.791197	1.144209
2000-01-10	0.030876	0.371900	0.061932

```
In [186]: tsdf.transform("abs")
```

```
Out[186]:
```

	A	B	C
2000-01-01	0.428759	0.864890	0.675341
2000-01-02	0.168731	1.338144	1.279321
2000-01-03	1.621034	0.438107	0.903794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	1.240447	0.201052
2000-01-09	0.157795	0.791197	1.144209
2000-01-10	0.030876	0.371900	0.061932

```
In [187]: tsdf.transform(lambda x: x.abs())
```

```
Out[187]:
```

	A	B	C
2000-01-01	0.428759	0.864890	0.675341
2000-01-02	0.168731	1.338144	1.279321
2000-01-03	1.621034	0.438107	0.903794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	1.240447	0.201052

```
2000-01-09  0.157795  0.791197  1.144209
2000-01-10  0.030876  0.371900  0.061932
```

Here `transform()` received a single function; this is equivalent to a `ufunc` application.

```
In [188]: np.abs(tsdf)
```

```
Out[188]:
```

	A	B	C
2000-01-01	0.428759	0.864890	0.675341
2000-01-02	0.168731	1.338144	1.279321
2000-01-03	1.621034	0.438107	0.903794
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	1.240447	0.201052
2000-01-09	0.157795	0.791197	1.144209
2000-01-10	0.030876	0.371900	0.061932

Passing a single function to `.transform()` with a `Series` will yield a single `Series` in return.

```
In [189]: tsdf["A"].transform(np.abs)
```

```
Out[189]:
```

```
2000-01-01    0.428759
2000-01-02    0.168731
2000-01-03    1.621034
2000-01-04         NaN
2000-01-05         NaN
2000-01-06         NaN
2000-01-07         NaN
2000-01-08    0.254374
2000-01-09    0.157795
2000-01-10    0.030876
Freq: D, Name: A, dtype: float64
```

Transform with multiple functions

Passing multiple functions will yield a column MultiIndexed DataFrame. The first level will be the original frame column names; the second level will be the names of the transforming functions.

```
In [190]: tsdf.transform([np.abs, lambda x: x + 1])
```

```
Out[190]:
```

	A		B		C	
	absolute	<lambda>	absolute	<lambda>	absolute	<lambda>
2000-01-01	0.428759	0.571241	0.864890	0.135110	0.675341	0.324659
2000-01-02	0.168731	0.831269	1.338144	2.338144	1.279321	-0.279321
2000-01-03	1.621034	-0.621034	0.438107	1.438107	0.903794	1.903794
2000-01-04	NaN	NaN	NaN	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN	NaN	NaN	NaN
2000-01-08	0.254374	1.254374	1.240447	-0.240447	0.201052	0.798948
2000-01-09	0.157795	0.842205	0.791197	1.791197	1.144209	-0.144209
2000-01-10	0.030876	0.969124	0.371900	1.371900	0.061932	1.061932

Passing multiple functions to a Series will yield a DataFrame. The resulting column names will be the transforming functions.

```
In [191]: tsdf["A"].transform([np.abs, lambda x: x + 1])
```

```
Out[191]:
```

	absolute	<lambda>
2000-01-01	0.428759	0.571241
2000-01-02	0.168731	0.831269
2000-01-03	1.621034	-0.621034
2000-01-04	NaN	NaN
2000-01-05	NaN	NaN
2000-01-06	NaN	NaN
2000-01-07	NaN	NaN
2000-01-08	0.254374	1.254374
2000-01-09	0.157795	0.842205
2000-01-10	0.030876	0.969124

Transforming with a dict

Passing a dict of functions will allow selective transforming per column.

```
In [192]: tsdf.transform({"A": np.abs, "B": lambda x: x + 1})
```

```
Out[192]:
```

	A	B
2000-01-01	0.428759	0.135110
2000-01-02	0.168731	2.338144
2000-01-03	1.621034	1.438107
2000-01-04	NaN	NaN
2000-01-05	NaN	NaN
2000-01-06	NaN	NaN
2000-01-07	NaN	NaN
2000-01-08	0.254374	-0.240447
2000-01-09	0.157795	1.791197
2000-01-10	0.030876	1.371900

Passing a dict of lists will generate a MultiIndexed DataFrame with these selective transforms.

```
In [193]: tsdf.transform({"A": np.abs, "B": [lambda x: x + 1, "sqrt"]})
Out[193]:
```

	A	B	
	absolute	<lambda>	sqrt
2000-01-01	0.428759	0.135110	NaN
2000-01-02	0.168731	2.338144	1.156782
2000-01-03	1.621034	1.438107	0.661897
2000-01-04	NaN	NaN	NaN
2000-01-05	NaN	NaN	NaN
2000-01-06	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN
2000-01-08	0.254374	-0.240447	NaN
2000-01-09	0.157795	1.791197	0.889493
2000-01-10	0.030876	1.371900	0.609836

Applying elementwise functions

Since not all functions can be vectorized (accept NumPy arrays and return another array or value), the methods `map()` on DataFrame and analogously `map()` on Series accept any Python function taking a single value and returning a single value. For example:

```

In [194]: df4 = df.copy()

In [195]: df4
Out[195]:
      one      two      three
a  1.394981  1.772517      NaN
b  0.343054  1.912123 -0.050390
c  0.695246  1.478369  1.227435
d      NaN  0.279344 -0.613172

In [196]: def f(x):
.....:     return len(str(x))
.....:

In [197]: df4["one"].map(f)
Out[197]:
a      18
b      19
c      18
d       3
Name: one, dtype: int64

In [198]: df4.map(f)
Out[198]:
      one  two  three
a     18   17     3
b     19   18    20
c     18   18    16
d      3   19    19

```

`Series.map()` has an additional feature; it can be used to easily “link” or “map” values defined by a secondary series. This is closely related to [merging/joining functionality](#):

```
In [199]: s = pd.Series(
.....:     ["six", "seven", "six", "seven", "six"], index=["a", "b", "c", "d", "e"]
.....: )
.....:

In [200]: t = pd.Series({"six": 6.0, "seven": 7.0})

In [201]: s
Out[201]:
a      six
b     seven
c      six
d     seven
e      six
dtype: str

In [202]: s.map(t)
Out[202]:
a      6.0
b      7.0
c      6.0
d      7.0
e      6.0
dtype: float64
```

Reindexing and altering labels

`reindex()` is the fundamental data alignment method in pandas. It is used to implement nearly all other features relying on label-alignment functionality. To *reindex* means to conform the data to match a given set of labels along a particular axis. This accomplishes several things:

- Reorders the existing data to match a new set of labels
- Inserts missing value (NA) markers in label locations where no data for that label existed

- If specified, **fill** data for missing labels using logic (highly relevant to working with time series data)

Here is a simple example:

```
In [203]: s = pd.Series(np.random.randn(5), index=["a", "b", "c", "d", "e"])
```

```
In [204]: s
```

```
Out[204]:
```

```
a    1.695148
b    1.328614
c    1.234686
d   -0.385845
e   -1.326508
dtype: float64
```

```
In [205]: s.reindex(["e", "b", "f", "d"])
```

```
Out[205]:
```

```
e   -1.326508
b    1.328614
f         NaN
d   -0.385845
dtype: float64
```

Here, the **f** label was not contained in the Series and hence appears as **NaN** in the result.

With a DataFrame, you can simultaneously reindex the index and columns:


```

In [206]: df
Out[206]:
      one      two      three
a  1.394981  1.772517      NaN
b  0.343054  1.912123 -0.050390
c  0.695246  1.478369  1.227435
d      NaN  0.279344 -0.613172

In [207]: df.reindex(index=["c", "f", "b"], columns=["three", "two", "one"])
Out[207]:
      three      two      one
c  1.227435  1.478369  0.695246
f      NaN      NaN      NaN
b -0.050390  1.912123  0.343054

```

Note that the `Index` objects containing the actual axis labels can be **shared** between objects. So if we have a Series and a DataFrame, the following can be done:

```

In [208]: rs = s.reindex(df.index)

In [209]: rs
Out[209]:
a    1.695148
b    1.328614
c    1.234686
d   -0.385845
dtype: float64

In [210]: rs.index is df.index
Out[210]: True

```

This means that the reindexed Series's index is the same Python object as the DataFrame's index.

`DataFrame.reindex()` also supports an “axis-style” calling convention, where you specify a single `labels` argument and the `axis` it applies to.

```
In [211]: df.reindex(["c", "f", "b"], axis="index")
```

```
Out[211]:
```

	one	two	three
c	0.695246	1.478369	1.227435
f	NaN	NaN	NaN
b	0.343054	1.912123	-0.050390

```
In [212]: df.reindex(["three", "two", "one"], axis="columns")
```

```
Out[212]:
```

	three	two	one
a	NaN	1.772517	1.394981
b	-0.050390	1.912123	0.343054
c	1.227435	1.478369	0.695246
d	-0.613172	0.279344	NaN

➡ See also

[MultiIndex / Advanced Indexing](#) is an even more concise way of doing reindexing.

i Note

When writing performance-sensitive code, there is a good reason to spend some time becoming a reindexing ninja: **many operations are faster on pre-aligned data**. Adding two unaligned DataFrames internally triggers a reindexing step. For exploratory analysis you will hardly notice the difference (because `reindex` has been heavily optimized), but when CPU cycles matter sprinkling a few explicit `reindex` calls here and there can have an impact.

Reindexing to align with another object

You may wish to take an object and reindex its axes to be labeled the same as another object. While the syntax for this is straightforward albeit verbose, it is a common enough operation that the

`reindex_like()` method is available to make this simpler:

```
In [213]: df2 = df.reindex(["a", "b", "c"], columns=["one", "two"])
```

```
In [214]: df3 = df2 - df2.mean()
```

```
In [215]: df2
```

```
Out[215]:
```

	one	two
a	1.394981	1.772517
b	0.343054	1.912123
c	0.695246	1.478369

```
In [216]: df3
```

```
Out[216]:
```

	one	two
a	0.583888	0.051514
b	-0.468040	0.191120
c	-0.115848	-0.242634

```
In [217]: df.reindex_like(df2)
```

```
Out[217]:
```

	one	two
a	1.394981	1.772517
b	0.343054	1.912123
c	0.695246	1.478369

Aligning objects with each other with `align`

The `align()` method is the fastest way to simultaneously align two objects. It supports a `join` argument (related to [joining and merging](#)):

- `join='outer'`: take the union of the indexes (default)
- `join='left'`: use the calling object's index
- `join='right'`: use the passed object's index
- `join='inner'`: intersect the indexes

It returns a tuple with both of the reindexed Series:

```
In [218]: s = pd.Series(np.random.randn(5), index=["a", "b", "c", "d", "e"])

In [219]: s1 = s[:4]

In [220]: s2 = s[1:]

In [221]: s1.align(s2)
Out[221]:
(a    -0.186646
 b    -1.692424
 c    -0.303893
 d    -1.425662
 e         NaN
dtype: float64,
 a         NaN
 b    -1.692424
 c    -0.303893
 d    -1.425662
 e     1.114285
dtype: float64)

In [222]: s1.align(s2, join="inner")
Out[222]:
(b    -1.692424
 c    -0.303893
 d    -1.425662
dtype: float64,
 b    -1.692424
 c    -0.303893
 d    -1.425662
dtype: float64)

In [223]: s1.align(s2, join="left")
Out[223]:
(a    -0.186646
 b    -1.692424
 c    -0.303893
 d    -1.425662
dtype: float64,
 a         NaN
```

```
b    -1.692424
c    -0.303893
d    -1.425662
dtype: float64)
```

For DataFrames, the join method will be applied to both the index and the columns by default:

```
In [224]: df.align(df2, join="inner")
Out[224]:
(
      one      two
a  1.394981  1.772517
b  0.343054  1.912123
c  0.695246  1.478369,
      one      two
a  1.394981  1.772517
b  0.343054  1.912123
c  0.695246  1.478369)
```

You can also pass an `axis` option to only align on the specified axis:

```
In [225]: df.align(df2, join="inner", axis=0)
Out[225]:
(
      one      two      three
a  1.394981  1.772517      NaN
b  0.343054  1.912123 -0.050390
c  0.695246  1.478369  1.227435,
      one      two
a  1.394981  1.772517
b  0.343054  1.912123
c  0.695246  1.478369)
```

If you pass a Series to `DataFrame.align()`, you can choose to align both objects either on the DataFrame's index or columns using the `axis` argument:

```
In [226]: df.align(df2.iloc[0], axis=1)
Out[226]:
(
      one      three      two
a  1.394981      NaN  1.772517
b  0.343054 -0.050390  1.912123
c  0.695246  1.227435  1.478369
d      NaN -0.613172  0.279344,
one      1.394981
three      NaN
two      1.772517
Name: a, dtype: float64)
```

Filling while reindexing

`reindex()` takes an optional parameter `method` which is a filling method chosen from the following table:

Method	Action
ffill	Fill values forward
bfill	Fill values backward
nearest	Fill from the nearest index value

We illustrate these fill methods on a simple Series:

```
In [227]: rng = pd.date_range("1/3/2000", periods=8)
```

```
In [228]: ts = pd.Series(np.random.randn(8), index=rng)
```

```
In [229]: ts2 = ts.iloc[[0, 3, 6]]
```

```
In [230]: ts
```

```
Out[230]:
```

```
2000-01-03    0.183051
2000-01-04    0.400528
2000-01-05   -0.015083
2000-01-06    2.395489
2000-01-07    1.414806
2000-01-08    0.118428
2000-01-09    0.733639
2000-01-10   -0.936077
Freq: D, dtype: float64
```

```
In [231]: ts2
```

```
Out[231]:
```

```
2000-01-03    0.183051
2000-01-06    2.395489
2000-01-09    0.733639
Freq: 3D, dtype: float64
```

```
In [232]: ts2.reindex(ts.index)
```

```
Out[232]:
```

```
2000-01-03    0.183051
2000-01-04         NaN
2000-01-05         NaN
2000-01-06    2.395489
2000-01-07         NaN
2000-01-08         NaN
2000-01-09    0.733639
2000-01-10         NaN
Freq: D, dtype: float64
```

```
In [233]: ts2.reindex(ts.index, method="ffill")
```

```
Out[233]:
```

```
2000-01-03    0.183051
```



```
2000-01-04    0.183051
2000-01-05    0.183051
2000-01-06    2.395489
2000-01-07    2.395489
2000-01-08    2.395489
2000-01-09    0.733639
2000-01-10    0.733639
Freq: D, dtype: float64
```

```
In [234]: ts2.reindex(ts.index, method="bfill")
```

```
Out[234]:
```

```
2000-01-03    0.183051
2000-01-04    2.395489
2000-01-05    2.395489
2000-01-06    2.395489
2000-01-07    0.733639
2000-01-08    0.733639
2000-01-09    0.733639
2000-01-10         NaN
Freq: D, dtype: float64
```

```
In [235]: ts2.reindex(ts.index, method="nearest")
```

```
Out[235]:
```

```
2000-01-03    0.183051
2000-01-04    0.183051
2000-01-05    2.395489
2000-01-06    2.395489
2000-01-07    2.395489
2000-01-08    0.733639
2000-01-09    0.733639
2000-01-10    0.733639
Freq: D, dtype: float64
```

These methods require that the indexes are **ordered** increasing or decreasing.

Note that the same result could have been achieved using [ffill](#) (except for `method='nearest'`) or [interpolate](#):

```
In [236]: ts2.reindex(ts.index).ffill()
Out[236]:
2000-01-03    0.183051
2000-01-04    0.183051
2000-01-05    0.183051
2000-01-06    2.395489
2000-01-07    2.395489
2000-01-08    2.395489
2000-01-09    0.733639
2000-01-10    0.733639
Freq: D, dtype: float64
```

`reindex()` will raise a `ValueError` if the index is not monotonically increasing or decreasing. `fillna()` and `interpolate()` will not perform any checks on the order of the index.

Limits on filling while reindexing

The `limit` and `tolerance` arguments provide additional control over filling while reindexing. `Limit` specifies the maximum count of consecutive matches:

```
In [237]: ts2.reindex(ts.index, method="ffill", limit=1)
Out[237]:
2000-01-03    0.183051
2000-01-04    0.183051
2000-01-05         NaN
2000-01-06    2.395489
2000-01-07    2.395489
2000-01-08         NaN
2000-01-09    0.733639
2000-01-10    0.733639
Freq: D, dtype: float64
```

In contrast, tolerance specifies the maximum distance between the index and indexer values:

```
In [238]: ts2.reindex(ts.index, method="ffill", tolerance="1 day")
Out[238]:
2000-01-03    0.183051
2000-01-04    0.183051
2000-01-05         NaN
2000-01-06    2.395489
2000-01-07    2.395489
2000-01-08         NaN
2000-01-09    0.733639
2000-01-10    0.733639
Freq: D, dtype: float64
```

Notice that when used on a `DatetimeIndex`, `TimedeltaIndex` or `PeriodIndex`, `tolerance` will be coerced into a `Timedelta` if possible. This allows you to specify tolerance with appropriate strings.

Dropping labels from an axis

A method closely related to `reindex` is the `drop()` function. It removes a set of labels from an axis:

```
In [239]: df
Out[239]:
```

	one	two	three
a	1.394981	1.772517	NaN
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [240]: df.drop(["a", "d"], axis=0)
Out[240]:
```

	one	two	three
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435

```
In [241]: df.drop(["one"], axis=1)
Out[241]:
```

	two	three
a	1.772517	NaN
b	1.912123	-0.050390
c	1.478369	1.227435
d	0.279344	-0.613172

Note that the following also works, but is a bit less obvious / clean:

```
In [242]: df.reindex(df.index.difference(["a", "d"]))
Out[242]:
```

	one	two	three
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435

Renaming / mapping labels

The `rename()` method allows you to relabel an axis based on some mapping (a dict or Series) or an arbitrary function.

```
In [243]: s
Out[243]:
a    -0.186646
b    -1.692424
c    -0.303893
d    -1.425662
e     1.114285
dtype: float64

In [244]: s.rename(str.upper)
Out[244]:
A    -0.186646
B    -1.692424
C    -0.303893
D    -1.425662
E     1.114285
dtype: float64
```

If you pass a function, it must return a value when called with any of the labels (and must produce a set of unique values). A dict or Series can also be used:

```
In [245]: df.rename(
.....:     columns={"one": "foo", "two": "bar"},
.....:     index={"a": "apple", "b": "banana", "d": "durian"},
.....: )
.....:
Out[245]:
```

	foo	bar	three
apple	1.394981	1.772517	NaN
banana	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
durian	NaN	0.279344	-0.613172

If the mapping doesn't include a column/index label, it isn't renamed. Note that extra labels in the mapping don't throw an error.

`DataFrame.rename()` also supports an "axis-style" calling convention, where you specify a single `mapper` and the `axis` to apply that mapping to.

```
In [246]: df.rename({"one": "foo", "two": "bar"}, axis="columns")
```

```
Out[246]:
```

	foo	bar	three
a	1.394981	1.772517	NaN
b	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
d	NaN	0.279344	-0.613172

```
In [247]: df.rename({"a": "apple", "b": "banana", "d": "durian"}, axis="index")
```

```
Out[247]:
```

	one	two	three
apple	1.394981	1.772517	NaN
banana	0.343054	1.912123	-0.050390
c	0.695246	1.478369	1.227435
durian	NaN	0.279344	-0.613172

Finally, `rename()` also accepts a scalar or list-like for altering the `Series.name` attribute.

```
In [248]: s.rename("scalar-name")
```

```
Out[248]:
```

a	-0.186646
b	-1.692424
c	-0.303893
d	-1.425662
e	1.114285

Name: scalar-name, dtype: float64

The methods `DataFrame.rename_axis()` and `Series.rename_axis()` allow specific names of a `MultiIndex` to be changed (as opposed to the labels).

```
In [249]: df = pd.DataFrame(
.....:     {"x": [1, 2, 3, 4, 5, 6], "y": [10, 20, 30, 40, 50, 60]},
.....:     index=pd.MultiIndex.from_product(
.....:         [["a", "b", "c"], [1, 2]], names=["let", "num"]
.....:     ),
.....: )
.....:
```

```
In [250]: df
```

```
Out[250]:
```

		x	y
let num			
a	1	1	10
	2	2	20
b	1	3	30
	2	4	40
c	1	5	50
	2	6	60

```
In [251]: df.rename_axis(index={"let": "abc"})
```

```
Out[251]:
```

		x	y
abc num			
a	1	1	10
	2	2	20
b	1	3	30
	2	4	40
c	1	5	50
	2	6	60

```
In [252]: df.rename_axis(index=str.upper)
```

```
Out[252]:
```

		x	y
LET NUM			
a	1	1	10
	2	2	20
b	1	3	30
	2	4	40


```
c    1    5  50
     2    6  60
```

Iteration

The behavior of basic iteration over pandas objects depends on the type. When iterating over a Series, it is regarded as array-like, and basic iteration produces the values. DataFrames follow the dict-like convention of iterating over the “keys” of the objects.

In short, basic iteration (`for i in object`) produces:

- **Series:** values
- **DataFrame:** column labels

Thus, for example, iterating over a DataFrame gives you the column names:

```
In [253]: df = pd.DataFrame(
.....:     {"col1": np.random.randn(3), "col2": np.random.randn(3)}, index=["a", "b", "c"]
.....: )
.....:

In [254]: for col in df:
.....:     print(col)
.....:
col1
col2
```

pandas objects also have the dict-like `items()` method to iterate over the (key, value) pairs.

To iterate over the rows of a DataFrame, you can use the following methods:

- `iterrows()`: Iterate over the rows of a DataFrame as (index, Series) pairs. This converts the rows to Series objects, which can change the dtypes and has some performance implications.
- `itertuples()`: Iterate over the rows of a DataFrame as namedtuples of the values. This is a lot faster than `iterrows()`, and is in most cases preferable to use to iterate over the values of a DataFrame.

Warning

Iterating through pandas objects is generally **slow**. In many cases, iterating manually over the rows is not needed and can be avoided with one of the following approaches:

- Look for a *vectorized* solution: many operations can be performed using built-in methods or NumPy functions, (boolean) indexing, ...
- When you have a function that cannot work on the full DataFrame/Series at once, it is better to use `apply()` instead of iterating over the values. See the docs on [function application](#).
- If you need to do iterative manipulations on the values but performance is important, consider writing the inner loop with cython or numba. See the [enhancing performance](#) section for some examples of this approach.

Warning

You should **never modify** something you are iterating over. This is not guaranteed to work in all cases. Depending on the data types, the iterator returns a copy and not a view, and writing to it will have no effect!

For example, in the following case setting the value has no effect:

```
In [255]: df = pd.DataFrame({"a": [1, 2, 3], "b": ["a", "b", "c"]})
```

```
In [256]: for index, row in df.iterrows():
.....:     row["a"] = 10
.....:
```

```
In [257]: df
```

```
Out[257]:
```

	a	b
0	1	a
1	2	b
2	3	c

items

Consistent with the dict-like interface, `items()` iterates through key-value pairs:

- **Series:** (index, scalar value) pairs
- **DataFrame:** (column, Series) pairs

For example:

```
In [258]: for label, ser in df.items():
.....:     print(label)
.....:     print(ser)
.....:
a
0    1
1    2
2    3
Name: a, dtype: int64
b
0    a
1    b
2    c
Name: b, dtype: str
```

iterrows

`iterrows()` allows you to iterate through the rows of a DataFrame as Series objects. It returns an iterator yielding each index value along with a Series containing the data in each row:

```
In [259]: for row_index, row in df.iterrows():
.....:     print(row_index, row, sep="\n")
.....:
0
a    1
b    a
Name: 0, dtype: object
1
a    2
b    b
Name: 1, dtype: object
2
a    3
b    c
Name: 2, dtype: object
```

Note

Because `iterrows()` returns a Series for each row, it does **not** preserve dtypes across the rows (dtypes are preserved across columns for DataFrames). For example,

```
In [260]: df_orig = pd.DataFrame([[1, 1.5]], columns=["int", "float"])
```

```
In [261]: df_orig.dtypes
```

```
Out[261]:
```

```
int      int64
```

```
float    float64
```

```
dtype: object
```

```
In [262]: row = next(df_orig.iterrows())[1]
```

```
In [263]: row
```

```
Out[263]:
```

```
int      1.0
```

```
float    1.5
```

```
Name: 0, dtype: float64
```

All values in `row`, returned as a Series, are now upcasted to floats, also the original integer value in column `x`:

```
In [264]: row["int"].dtype
```

```
Out[264]: dtype('float64')
```

```
In [265]: df_orig["int"].dtype
```

```
Out[265]: dtype('int64')
```

To preserve dtypes while iterating over the rows, it is better to use `itertuples()` which returns namedtuples of the values and which is generally much faster than `iterrows()`.

For instance, a contrived way to transpose the DataFrame would be:

```
In [266]: df2 = pd.DataFrame({"x": [1, 2, 3], "y": [4, 5, 6]})

In [267]: print(df2)
   x  y
0  1  4
1  2  5
2  3  6

In [268]: print(df2.T)
   0  1  2
x  1  2  3
y  4  5  6

In [269]: df2_t = pd.DataFrame({idx: values for idx, values in df2.iterrows()})

In [270]: print(df2_t)
   0  1  2
x  1  2  3
y  4  5  6
```

itertuples

The `itertuples()` method will return an iterator yielding a namedtuple for each row in the DataFrame. The first element of the tuple will be the row's corresponding index value, while the remaining values are the row values.

For instance:

```
In [271]: for row in df.itertuples():
.....:     print(row)
.....:
Pandas(Index=0, a=1, b='a')
Pandas(Index=1, a=2, b='b')
Pandas(Index=2, a=3, b='c')
```

This method does not convert the row to a Series object; it merely returns the values inside a namedtuple. Therefore, `itertuples()` preserves the data type of the values and is generally faster than `iterrows()`.

Note

The column names will be renamed to positional names if they are invalid Python identifiers, repeated, or start with an underscore. With a large number of columns (>255), regular tuples are returned.

.dt accessor

`Series` has an accessor to succinctly return datetime like properties for the *values* of the Series, if it is a datetime/period like Series. This will return a Series, indexed like the existing Series.


```
# datetime
In [272]: s = pd.Series(pd.date_range("20130101 09:10:12", periods=4))

In [273]: s
Out[273]:
0    2013-01-01 09:10:12
1    2013-01-02 09:10:12
2    2013-01-03 09:10:12
3    2013-01-04 09:10:12
dtype: datetime64[us]

In [274]: s.dt.hour
Out[274]:
0     9
1     9
2     9
3     9
dtype: int32

In [275]: s.dt.second
Out[275]:
0    12
1    12
2    12
3    12
dtype: int32

In [276]: s.dt.day
Out[276]:
0     1
1     2
2     3
3     4
dtype: int32
```

This enables nice expressions like this:

```
In [277]: s[s.dt.day == 2]
Out[277]:
1    2013-01-02 09:10:12
dtype: datetime64[us]
```

You can easily produces tz aware transformations:

```
In [278]: stz = s.dt.tz_localize("US/Eastern")

In [279]: stz
Out[279]:
0    2013-01-01 09:10:12-05:00
1    2013-01-02 09:10:12-05:00
2    2013-01-03 09:10:12-05:00
3    2013-01-04 09:10:12-05:00
dtype: datetime64[us, US/Eastern]

In [280]: stz.dt.tz
Out[280]: zoneinfo.ZoneInfo(key='US/Eastern')
```

You can also chain these types of operations:

```
In [281]: s.dt.tz_localize("UTC").dt.tz_convert("US/Eastern")
Out[281]:
0    2013-01-01 04:10:12-05:00
1    2013-01-02 04:10:12-05:00
2    2013-01-03 04:10:12-05:00
3    2013-01-04 04:10:12-05:00
dtype: datetime64[us, US/Eastern]
```

You can also format datetime values as strings with `Series.dt.strftime()` which supports the same format as the standard `strftime()`.

```
# DatetimeIndex
In [282]: s = pd.Series(pd.date_range("20130101", periods=4))

In [283]: s
Out[283]:
0    2013-01-01
1    2013-01-02
2    2013-01-03
3    2013-01-04
dtype: datetime64[us]

In [284]: s.dt.strftime("%Y/%m/%d")
Out[284]:
0    2013/01/01
1    2013/01/02
2    2013/01/03
3    2013/01/04
dtype: str
```

```
# PeriodIndex
In [285]: s = pd.Series(pd.period_range("20130101", periods=4))

In [286]: s
Out[286]:
0    2013-01-01
1    2013-01-02
2    2013-01-03
3    2013-01-04
dtype: period[D]

In [287]: s.dt.strftime("%Y/%m/%d")
Out[287]:
0    2013/01/01
1    2013/01/02
2    2013/01/03
3    2013/01/04
dtype: str
```

The `.dt` accessor works for period and timedelta dtypes.

```
# period
In [288]: s = pd.Series(pd.period_range("20130101", periods=4, freq="D"))

In [289]: s
Out[289]:
0    2013-01-01
1    2013-01-02
2    2013-01-03
3    2013-01-04
dtype: period[D]

In [290]: s.dt.year
Out[290]:
0    2013
1    2013
2    2013
3    2013
dtype: int64

In [291]: s.dt.day
Out[291]:
0    1
1    2
2    3
3    4
dtype: int64
```

```
# timedelta
In [292]: s = pd.Series(pd.timedelta_range("1 day 00:00:05", periods=4, freq="s"))
```

```
In [293]: s
```

```
Out[293]:
```

```
0    1 days 00:00:05
1    1 days 00:00:06
2    1 days 00:00:07
3    1 days 00:00:08
dtype: timedelta64[us]
```

```
In [294]: s.dt.days
```

```
Out[294]:
```

```
0     1
1     1
2     1
3     1
dtype: int64
```

```
In [295]: s.dt.seconds
```

```
Out[295]:
```

```
0     5
1     6
2     7
3     8
dtype: int32
```

```
In [296]: s.dt.components
```

```
Out[296]:
```

	days	hours	minutes	seconds	milliseconds	microseconds	nanoseconds
0	1	0	0	5	0	0	0
1	1	0	0	6	0	0	0
2	1	0	0	7	0	0	0
3	1	0	0	8	0	0	0

Note

`Series.dt` will raise a `TypeError` if you access with a non-datetime-like values.

Vectorized string methods

Series is equipped with a set of string processing methods that make it easy to operate on each element of the array. Perhaps most importantly, these methods exclude missing/NA values automatically. These are accessed via the Series's `str` attribute and generally have names matching the equivalent (scalar) built-in string methods. For example:

```
In [297]: s = pd.Series(
.....:     ["A", "B", "C", "Aaba", "Baca", np.nan, "CABA", "dog", "cat"], dtype="string"
.....: )
.....:

In [298]: s.str.lower()
Out[298]:
0      a
1      b
2      c
3    aaba
4    baca
5    <NA>
6    caba
7    dog
8    cat
dtype: string
```

Powerful pattern-matching methods are provided as well, but note that pattern-matching generally uses [regular expressions](#) by default (and in some cases always uses them).

Note

Prior to pandas 1.0, string methods were only available on `object`-dtype `Series`. pandas 1.0 added the `StringDtype` which is dedicated to strings. See [Text data types](#) for more.

Please see [Vectorized String Methods](#) for a complete description.

Sorting

pandas supports three kinds of sorting: sorting by index labels, sorting by column values, and sorting by a combination of both.

By index

The `Series.sort_index()` and `DataFrame.sort_index()` methods are used to sort a pandas object by its index levels.

```
In [299]: df = pd.DataFrame(
.....:     {
.....:         "one": pd.Series(np.random.randn(3), index=["a", "b", "c"]),
.....:         "two": pd.Series(np.random.randn(4), index=["a", "b", "c", "d"]),
.....:         "three": pd.Series(np.random.randn(3), index=["b", "c", "d"]),
.....:     }
.....: )
.....:
```

```
In [300]: unsorted_df = df.reindex(
.....:     index=["a", "d", "c", "b"], columns=["three", "two", "one"]
.....: )
.....:
```

```
In [301]: unsorted_df
```

```
Out[301]:
```

	three	two	one
a	NaN	-1.152244	0.562973
d	-0.252916	-0.109597	NaN
c	1.273388	-0.167123	0.640382
b	-0.098217	0.009797	-1.299504

```
# DataFrame
```

```
In [302]: unsorted_df.sort_index()
```

```
Out[302]:
```

	three	two	one
a	NaN	-1.152244	0.562973
b	-0.098217	0.009797	-1.299504
c	1.273388	-0.167123	0.640382
d	-0.252916	-0.109597	NaN

```
In [303]: unsorted_df.sort_index(ascending=False)
```

```
Out[303]:
```

	three	two	one
d	-0.252916	-0.109597	NaN
c	1.273388	-0.167123	0.640382
b	-0.098217	0.009797	-1.299504
a	NaN	-1.152244	0.562973

```
In [304]: unsorted_df.sort_index(axis=1)
```



```

Out[304]:
      one      three      two
a  0.562973      NaN -1.152244
d      NaN -0.252916 -0.109597
c  0.640382  1.273388 -0.167123
b -1.299504 -0.098217  0.009797

# Series
In [305]: unsorted_df["three"].sort_index()
Out[305]:
a      NaN
b  -0.098217
c   1.273388
d  -0.252916
Name: three, dtype: float64

```

Sorting by index also supports a `key` parameter that takes a callable function to apply to the index being sorted. For `MultiIndex` objects, the key is applied per-level to the levels specified by `level`.

```

In [306]: s1 = pd.DataFrame({"a": ["B", "a", "C"], "b": [1, 2, 3], "c": [2, 3, 4]}).set_index(
.....:     list("ab")
.....: )
.....:

In [307]: s1
Out[307]:
      c
a b
B 1  2
a 2  3
C 3  4

```

```

In [308]: s1.sort_index(level="a")
Out[308]:
      c
a b
B 1  2
C 3  4
a 2  3

In [309]: s1.sort_index(level="a", key=lambda idx: idx.str.lower())
Out[309]:
      c
a b
a 2  3
B 1  2
C 3  4

```

For information on key sorting by value, see [value sorting](#).

By values

The `Series.sort_values()` method is used to sort a `Series` by its values. The `DataFrame.sort_values()` method is used to sort a `DataFrame` by its column or row values. The optional `by` parameter to `DataFrame.sort_values()` may be used to specify one or more columns to use to determine the sorted order.

```
In [310]: df1 = pd.DataFrame(
.....:     {"one": [2, 1, 1, 1], "two": [1, 3, 2, 4], "three": [5, 4, 3, 2]}
.....: )
.....:
```

```
In [311]: df1.sort_values(by="two")
```

```
Out[311]:
```

	one	two	three
0	2	1	5
2	1	2	3
1	1	3	4
3	1	4	2

The `by` parameter can take a list of column names, e.g.:

```
In [312]: df1[["one", "two", "three"]].sort_values(by=["one", "two"])
```

```
Out[312]:
```

	one	two	three
2	1	2	3
1	1	3	4
3	1	4	2
0	2	1	5

These methods have special treatment of NA values via the `na_position` argument:

```
In [313]: s[2] = np.nan
```

```
In [314]: s.sort_values()
```

```
Out[314]:
```

```
0      A
3    Aaba
1      B
4    Baca
6    CABA
8     cat
7     dog
2    <NA>
5    <NA>
dtype: string
```

```
In [315]: s.sort_values(na_position="first")
```

```
Out[315]:
```

```
2    <NA>
5    <NA>
0      A
3    Aaba
1      B
4    Baca
6    CABA
8     cat
7     dog
dtype: string
```

Sorting also supports a `key` parameter that takes a callable function to apply to the values being sorted.

```
In [316]: s1 = pd.Series(["B", "a", "C"])
```

```
In [317]: s1.sort_values()
```

```
Out[317]:
```

```
0    B
2    C
1    a
dtype: str
```

```
In [318]: s1.sort_values(key=lambda x: x.str.lower())
```

```
Out[318]:
```

```
1    a
0    B
2    C
dtype: str
```

`key` will be given the `Series` of values and should return a `Series` or array of the same shape with the transformed values. For `DataFrame` objects, the key is applied per column, so the key should still expect a `Series` and return a `Series`, e.g.

```
In [319]: df = pd.DataFrame({"a": ["B", "a", "C"], "b": [1, 2, 3]})
```

```
In [320]: df.sort_values(by="a")
```

```
Out[320]:
```

```
   a  b
0  B  1
2  C  3
1  a  2
```

```
In [321]: df.sort_values(by="a", key=lambda col: col.str.lower())
```

```
Out[321]:
```

```
   a  b
1  a  2
0  B  1
2  C  3
```

The name or type of each column can be used to apply different functions to different columns.

By indexes and values

Strings passed as the `by` parameter to `DataFrame.sort_values()` may refer to either columns or index level names.

```
# Build MultiIndex
In [322]: idx = pd.MultiIndex.from_tuples(
.....:     [("a", 1), ("a", 2), ("a", 2), ("b", 2), ("b", 1), ("b", 1)]
.....: )
.....:

In [323]: idx.names = ["first", "second"]

# Build DataFrame
In [324]: df_multi = pd.DataFrame({"A": np.arange(6, 0, -1)}, index=idx)

In [325]: df_multi
Out[325]:
```

		A
first	second	
a	1	6
	2	5
	2	4
b	2	3
	1	2
	1	1

Sort by 'second' (index) and 'A' (column)

```
In [326]: df_multi.sort_values(by=["second", "A"])
```

```
Out[326]:
```

		A
first	second	
b	1	1
	1	2
a	1	6
b	2	3
a	2	4
	2	5

Note

If a string matches both a column name and an index level name then a warning is issued and the column takes precedence. This will result in an ambiguity error in a future version.

searchsorted

Series has the `searchsorted()` method, which works similarly to `numpy.ndarray.searchsorted()`.

```
In [327]: ser = pd.Series([1, 2, 3])

In [328]: ser.searchsorted([0, 3])
Out[328]: array([0, 2])

In [329]: ser.searchsorted([0, 4])
Out[329]: array([0, 3])

In [330]: ser.searchsorted([1, 3], side="right")
Out[330]: array([1, 3])

In [331]: ser.searchsorted([1, 3], side="left")
Out[331]: array([0, 2])

In [332]: ser = pd.Series([3, 1, 2])

In [333]: ser.searchsorted([0, 3], sorter=np.argsort(ser))
Out[333]: array([0, 2])
```

smallest / largest values

`Series` has the `nsmallest()` and `nlargest()` methods which return the smallest or largest n values. For a large `Series` this can be much faster than sorting the entire Series and calling `head(n)` on the result.


```
In [334]: s = pd.Series(np.random.permutation(10))
```

```
In [335]: s
```

```
Out[335]:
```

```
0    2
1    0
2    3
3    7
4    1
5    5
6    9
7    6
8    8
9    4
```

```
dtype: int64
```

```
In [336]: s.sort_values()
```

```
Out[336]:
```

```
1    0
4    1
0    2
2    3
9    4
5    5
7    6
3    7
8    8
6    9
```

```
dtype: int64
```

```
In [337]: s.nsmallest(3)
```

```
Out[337]:
```

```
1    0
4    1
0    2
```

```
dtype: int64
```

```
In [338]: s.nlargest(3)
```

```
Out[338]:
```

```
6    9
```

```
8    8
3    7
dtype: int64
```

`DataFrame` also has the `nlargest` and `nsmallest` methods.

```
In [339]: df = pd.DataFrame(
.....:     {
.....:         "a": [-2, -1, 1, 10, 8, 11, -1],
.....:         "b": list("abdceff"),
.....:         "c": [1.0, 2.0, 4.0, 3.2, np.nan, 3.0, 4.0],
.....:     }
.....: )
.....:
```

```
In [340]: df.nlargest(3, "a")
```

```
Out[340]:
```

	a	b	c
5	11	f	3.0
3	10	c	3.2
4	8	e	NaN

```
In [341]: df.nlargest(5, ["a", "c"])
```

```
Out[341]:
```

	a	b	c
5	11	f	3.0
3	10	c	3.2
4	8	e	NaN
2	1	d	4.0
6	-1	f	4.0

```
In [342]: df.nsmallest(3, "a")
```

```
Out[342]:
```

	a	b	c
0	-2	a	1.0
1	-1	b	2.0
6	-1	f	4.0

```
In [343]: df.nsmallest(5, ["a", "c"])
```

```
Out[343]:
```

	a	b	c
0	-2	a	1.0
1	-1	b	2.0
6	-1	f	4.0

```
2  1  d  4.0
4  8  e  NaN
```

Sorting by a MultiIndex column

You must be explicit about sorting when the column is a MultiIndex, and fully specify all levels to `by`.

```
In [344]: df1.columns = pd.MultiIndex.from_tuples(
.....:     [("a", "one"), ("a", "two"), ("b", "three")]
.....: )
.....:
```

```
In [345]: df1.sort_values(by=("a", "two"))
```

```
Out[345]:
```

	a		b
	one	two	three
0	2	1	5
2	1	2	3
1	1	3	4
3	1	4	2

Copying

The `copy()` method on pandas objects copies the underlying data (though not the axis indexes, since they are immutable) and returns a new object. Note that **it is seldom necessary to copy objects**. For example, there are only a handful of ways to alter a DataFrame *in-place*:

- Inserting, deleting, or modifying a column.
- Assigning to the `index` or `columns` attributes.

- For homogeneous data, directly modifying the values via the `values` attribute or advanced indexing.

To be clear, no pandas method has the side effect of modifying your data; almost every method returns a new object, leaving the original object untouched. If the data is modified, it is because you did so explicitly.

dtypes

For the most part, pandas uses NumPy arrays and dtypes for Series or individual columns of a DataFrame. NumPy provides support for `float`, `int`, `bool`, `timedelta64[ns]` and `datetime64[ns]` (note that NumPy does not support timezone-aware datetimes).

pandas and third-party libraries *extend* NumPy's type system in a few places. This section describes the extensions pandas has made internally. See [Extension types](#) for how to write your own extension that works with pandas. See [the ecosystem page](#) for a list of third-party libraries that have implemented an extension.

The following table lists all of pandas extension types. For methods requiring `dtype` arguments, strings can be specified as indicated. See the respective documentation sections for more on each type.

Kind of Data	Data Type	Scalar	Array	String Aliases
tz-aware datetime	<code>DatetimeTZDtype</code>	<code>Timestamp</code>	<code>arrays.DatetimeArray</code>	<code>'datetime64[ns, <tz>]'</code>
Categorical	<code>CategoricalDtype</code>	(none)	<code>Categorical</code>	<code>'category'</code>

Kind of Data	Data Type	Scalar	Array	String Aliases
period (time spans)	PeriodDtype	Period	arrays.PeriodArray 'Period[<freq>']	'period[<freq>']
sparse	SparseDtype	(none)	arrays.SparseArray	'Sparse', 'Sparse[int]', 'Sparse[float]'
intervals	IntervalDtype	Interval	arrays.IntervalArray	'interval', 'Interval', 'Interval[<numpy_dtype>]', 'Interval[datetime64[ns, <tz>]]', 'Interval[timedelta64[<freq>]
nullable integer	Int64Dtype , ...	(none)	arrays.IntegerArray	'Int8', 'Int16', 'Int32', 'Int64', 'UInt8', 'UInt16', 'UInt32', 'UInt64'
nullable float	Float64Dtype , ...	(none)	arrays.FloatingArray	'Float32', 'Float64'
Strings	StringDtype	str	arrays.StringArray	'string'
Boolean (with NA)	BooleanDtype	bool	arrays.BooleanArray	'boolean'

pandas has two ways to store strings.

1. `object` dtype, which can hold any Python object, including strings.
2. `StringDtype`, which is dedicated to strings.

Generally, we recommend using `StringDtype`. See [Text data types](#) for more.

Finally, arbitrary objects may be stored using the `object` dtype, but should be avoided to the extent possible (for performance and interoperability with other libraries and methods. See [object conversion](#)).

A convenient `dtypes` attribute for DataFrame returns a Series with the data type of each column.

```
In [346]: dft = pd.DataFrame(
.....:     {
.....:         "A": np.random.rand(3),
.....:         "B": 1,
.....:         "C": "foo",
.....:         "D": pd.Timestamp("20010102"),
.....:         "E": pd.Series([1.0] * 3).astype("float32"),
.....:         "F": False,
.....:         "G": pd.Series([1] * 3, dtype="int8"),
.....:     }
.....: )
.....:
```

```
In [347]: dft
```

```
Out[347]:
```

	A	B	C	D	E	F	G
0	0.035962	1	foo	2001-01-02	1.0	False	1
1	0.701379	1	foo	2001-01-02	1.0	False	1
2	0.281885	1	foo	2001-01-02	1.0	False	1

```
In [348]: dft.dtypes
```

```
Out[348]:
```

```
A          float64
B          int64
C           str
D    datetime64[us]
E          float32
F           bool
G           int8
dtype: object
```

On a `Series` object, use the `dtype` attribute.

```
In [349]: dft["A"].dtype
Out[349]: dtype('float64')
```


If a pandas object contains data with multiple dtypes *in a single column*, the dtype of the column will be chosen to accommodate all of the data types (`object` is the most general).

```
# these ints are coerced to floats
In [350]: pd.Series([1, 2, 3, 4, 5, 6.0])
Out[350]:
0    1.0
1    2.0
2    3.0
3    4.0
4    5.0
5    6.0
dtype: float64
```

```
# string data forces an ``object`` dtype
In [351]: pd.Series([1, 2, 3, 6.0, "foo"])
Out[351]:
0     1
1     2
2     3
3    6.0
4    foo
dtype: object
```

The number of columns of each type in a `DataFrame` can be found by calling `DataFrame.dtypes.value_counts()`.

```
In [352]: dft.dtypes.value_counts()
Out[352]:
float64      1
int64        1
str          1
datetime64[us] 1
float32      1
bool         1
int8         1
Name: count, dtype: int64
```

Numeric dtypes will propagate and can coexist in DataFrames. If a dtype is passed (either directly via the `dtype` keyword, a passed `ndarray`, or a passed `Series`), then it will be preserved in DataFrame operations. Furthermore, different numeric dtypes will **NOT** be combined. The following example will give you a taste.